

The Installed Veronica.v4 Foundation: Huihui-Qwen3-30B-A3B-Instruct-2507-abliterated

Document type: cited research paper (not model qualification, not legal advice, not a selection decision)

Subject checkpoint: Candidate A, revision e2f73ec7e99ee316beb8069ca90e4c3cbef8aa0f

Project: Veronica.v4 (C:\Users\raine\DEVELOPMENT\Projects\Veronica.v4)

Date: 2026-09-08

Status: the public alias Veronica is a wrapper rename. The weights on disk and on the RunPod volume are this Huihui derivative of Qwen3-30B-A3B-Instruct-2507. They are not Candidate B (Qwen3.8-27B).

Selection status remains benchmark_required.

Abstract

The model currently installed for Veronica.v4 text serving is a community abliterated derivative of Alibaba's Qwen3-30B-A3B-Instruct-2507, published by huihui-ai at an immutable Hugging Face revision pinned by this project. It is a 30.5 billion parameter Mixture-of-Experts causal LM with 3.3 billion activated parameters, 128 experts and 8 experts per token. The official base is Instruct-only and non-thinking: it does not emit think-tag blocks. Huihui's card states that refusals were reduced with a crude, proof-of-concept abliteration method, citing Sumandora's Hugging Face Transformers implementation of the Ardit et al. refusal-direction technique, plus an unspecified "new and faster method." This project has not fine-tuned, quantized, or LoRA-adapted those weights. It has hash-verified all 27 repository files (61,084,222,203 bytes), served them with vLLM 0.11.0 under the alias Veronica on an A100 80 GB, wrapped them with a local persona/mode prompt layer, and recorded first-chat and later smoke failures that block qualification. Candidate B (Huihui-Qwen3.8-27B-abliterated) is a different architecture and is not present on the model volume.

1. Identity of the installed checkpoint

Field	Installed value	Source
Public API alias	Veronica	config/model-registry.json; wrapper /v1/models
Hugging Face repository	huihui-ai/Huihui-Qwen3-30B-A3B-Instruct-2507-abliterated	Hub card; registry; validated manifest
Immutable revision	e2f73ec7e99ee316beb8069ca90e4c3cbef8aa0f	Hub metadata sha; registry; every successful start run
Hub lastModified	2025-08-02T05:55:21.000Z	provenance/candidate/hub-metadata.json
Declared license	Apache-2.0 (LICENSE file present; license_link points at the Qwen control LICENSE)	Hub card YAML; runs/2026-09-01-t2-qualification/license-review.md
Declared base	Qwen/Qwen3-30B-A3B-Instruct-2507	Hub base_model; card prose
Control revision (official, not served)	0d7cf23991f47feeb3a57ecb4c9cee8ea4a17bfe	registry; config/runpod-core.json
Architecture	Qwen3MoeForCausalLM / model_type: qwen3_moe	installed config.json
Parameters (Hub safetensors census)	30,532,122,624 BF16 tensors	Hub metadata safetensors.parameters.BF16

Field	Installed value	Source
Official card parameter claim	30.5B total, 3.3B activated, 29.9B non-embedding	Qwen Instruct-2507 card
Experts	128 total, 8 activated per token, 48 layers, GQA 32/4, hidden size 2048	installed config.json; Qwen card
Native context	262,144 tokens (Qwen card). This project currently serves --max-model-len 8192.	Qwen card; server-command.json
On-disk verified size	61,084,222,203 bytes, 27 files, 13 weight shards	validated-model-manifest.json
Volume path	/workspace/veronica-core/models/Huihui-Qwen3-30B-A3B-Instruct-2507-abliterated/e2f73ec7e99ee316beb8069ca90e4c3cbef8aa0f	validated manifest; server command
Network volume	v53gj9flzs, 300 GB, EUR-IS-1	config/runpod-core.json; volume inspection

The 2026-09-08 start run's validated file list, byte counts, and SHA-256 hashes are identical to the 2026-08-30 first-chat validated manifest. Parsed config.json from the T2 metadata snapshot equals the served config; only pretty-print whitespace differs (23f9d3f5... vs 0736f3ba... as raw SHA-256). That is a snapshot-encoding difference, not a different model.

Integrity check nuance: Hub expected-model-manifest.json carries LFS SHA-256 for the 13 weight shards plus tokenizer.json (14 files). The other 13 files are expected by size and git-blob SHA-1. prepare_runpod_model.py verify_files() hashes every file with SHA-256 and git-blob SHA-1, compares against whichever expected checksum exists, then writes SHA-256 for all 27 files into validated-model-manifest.json. So the weights are SHA-256-verified against Hub LFS; small text/config files are git-blob-verified, then locally SHA-256-recorded.

The official Qwen Instruct-2507 Hub dump packs the same BF16 parameter count as 16 shards (~61.06 GB mapped weights). Huihui's derivative is 13 shards totaling 61,084,222,203 bytes. Same architecture fields and matching generation_config.json (temperature 0.7, top_k 20, top_p 0.8); different shard layout, so shard-for-shard hash comparison against the official control is not valid. Compare whole-model provenance, not individual shard names.

2. What this model is not

Not Candidate B. Candidate B is huihui-ai/Huihui-Qwen3.8-27B-abliterated at 739e3c5b89849f6c238ce1e5b70008612ae42cdd, architecture Qwen3_5ForConditionalGeneration / qwen3_5, a dense multimodal model with native thinking and reasoning_effort. Its config.json SHA-256 (191e0af2...) does not match the installed config. The 2026-09-08 volume listing under /workspace/veronica-core/models/ contains only Huihui-Qwen3-30B-A3B-Instruct-2507-abliterated. T2 still requires transferring and hash-validating Candidate B and both official controls before any fair comparison.

Not a Veronica-trained foundation. The string Veronica is the stable public alias. vLLM is started with --served-model-name Veronica. The wrapper rewrites the model field on completions. No project run claims a training job, LoRA, QLoRA, or merge on these weights.

Not the Qwen3-30B-A3B thinking variant. Qwen published a separate Thinking-2507 sibling. Huihui also published Huihui-Qwen3-30B-A3B-Thinking-2507-abliterated. The installed repo name, card, and config.json are the Instruct-2507 line. The official Instruct card states that the model supports only non-thinking mode and does not generate think-tag blocks; enable_thinking=False is no longer required.

Not a quantized GGUF/EXL3/MLX serving copy. Third-party quants exist (for example mradermacher GGUF). This project serves BF16 safetensors with vLLM, dtype bfloat16.

3. Official Qwen3-30B-A3B-Instruct-2507

Qwen3-Instruct-2507 is the July 2025 update of the Qwen3-30B-A3B non-thinking mode (GitHub QwenLM/Qwen3 changelog: Instruct-2507 released 2025-07-30). The official card lists:

- Causal LM, pretraining plus post-training.
- 30.5B total / 3.3B activated MoE; 48 layers; 128 experts / 8 active; GQA 32 query / 4 KV heads.
- Native context 262,144 tokens. Optional Dual Chunk Attention plus MInference path toward ~1M tokens, which the card says needs on the order of 240 GB total GPU memory and is not what this project runs.
- Claimed gains versus the earlier non-thinking 30B-A3B on instruction following, reasoning, multilingual coverage, open-ended alignment, and long context.
- Reported official scores on the card include MMLU-Pro 78.4, AIME25 61.3, LiveCodeBench v6 43.2, IFEval 84.7, Arena-Hard v2 69.0, Creative Writing v3 86.0, BFCL-v3 65.1. Those numbers describe the official untouched Instruct-2507, not the ablated derivative, and not this project's 8,192-token vLLM 0.11.0 smoke runtime.
- Deployment note: transformers>=4.51.0 (older versions raise KeyError: 'qwen3_moe'). Suggested serving: vllm>=0.8.5 or sglang>=0.4.6.post1. Suggested sampling: temperature 0.7, top_p 0.8, top_k 20.
- Technical report: Qwen Team, *Qwen3 Technical Report*, arXiv:2505.09388.

The official control in this repo is pinned at 0d7cf23991f47feeb3a57ecb4c9cee8ea4a17bfe. That control is not downloaded as weights on the volume. Only its README, LICENSE, and related metadata snapshots were captured for T2.

4. Abliteration: what was changed in the weights

4.1 Scientific basis

Arditi, Obeso, Syed, Paleka, Panickssery, Gurnee, and Nanda (*Refusal in Language Models Is Mediated by a Single Direction*, arXiv:2406.11717, NeurIPS 2024) showed, across 13 open chat models up to 72B, that refusal is mediated by a one-dimensional residual-stream direction: erasing it reduces refusal of harmful instructions; adding it can induce refusal of harmless ones. They describe this as a white-box intervention with limited collateral on other capabilities in their study. That paper does not evaluate this Huihui checkpoint.

4.2 Implementation Huihui cites

The Huihui card says the uncensored derivative was "created with ablation (see remove-refusals-with-transformers)" and repeats Sumandora's own wording: "a crude, proof-of-concept implementation to remove refusals from an LLM model without using TransformerLens." Sumandora's repository:

- Computes a refusal direction from harmful vs harmless hidden-state means (compute_refusal_dir.py).

- At inference, inserts `AblationDecoderLayer` modules that subtract that direction's projection from activations (`direction_ablation_hook`). The public Sumandora tree has no weight-save / orthogonalization exporter.
- Credits Ardit et al. as the technique source.
- Warns that some Qwen implementations do not expose `model.model.layers`.

Huihui uploaded a full safetensor dump, not a runtime hook. The card does not say whether they baked an Ardit-style rank-one edit into `w_out` (embedding / attention-out / MLP-out) or used another unpublished converter. Hugging Face's `base_model:finetune` tag is Hub taxonomy for a derived repo, not proof that Huihui ran SFT, LoRA, or DPO.

Community tooling (Orion-zhen / Undi95 `abliteration`) writes the usual rank-1 update that subtracts a scaled outer product of the refusal direction from the weight matrix, and notes explicitly that `abliteration` is not full uncensorship: the model may still fail to refuse without becoming a different, more capable model.

4.3 What Huihui actually discloses for this repo

Pinned Candidate A README (revision `e2f73ec7...`):

1. Base = Qwen/Qwen3-30B-A3B-Instruct-2507.
2. Method = `abliteration` via the Sumandora recipe, described as crude PoC.
3. Extra sentence: "Ablation was performed using a new and faster method, which yields better results."
4. Usage warnings: reduced safety filtering; not for all audiences; user bears legal/ethical responsibility; recommended for research/testing rather than production; huihui.ai disclaims consequences.
5. Example inference script still passes `enable_thinking = not nothink` into `apply_chat_template`, leftover from thinking-capable siblings. The official Instruct-2507 card says thinking mode is not supported. This project's SOURCE-OF-TRUTH records the same limitation.

Huihui does not publish, for Candidate A:

- which layers were modified (Candidate B's card **does** say layers 18-51);
- the refusal-direction vector, scale α , or whether the change is activation-hook vs baked weights;
- a capability regression table against Instruct-2507;
- independent third-party audit of residual refusal or capability damage.

This project's T2 license/lineage review records those statements as repository-author declarations, not independently proven facts, and does not treat them as capability evidence.

What Veronica.v4 has done to the weights: nothing after that community ablation. Project principles forbid changing foundation weights during the alias/persona-wrapper stage. Bootstrap verifies hashes and atomically promotes a directory; it does not train.

5. License and redistribution (documented, not legal advice)

All four T2 identities (A, A-control, B, B-control) declare Apache-2.0 at their pinned revisions; each listing contains `LICENSE` and none listed `NOTICE`. Apache-2.0, subject to its terms, permits reproduction, derivative works, display, sublicense, and distribution, and includes a patent

license from contributors. Redistribution requires preserving the license, marking modified files, and retaining applicable notices. It does not grant trademarks and includes warranty/liability disclaimers.

Huihui's card `license_link` points at the Qwen Instruct-2507 LICENSE, which is consistent with a derivative that inherits the base license. The T2 review states that repository declarations do not replace final legal review, do not prove uploader authority or training-data rights, and must be re-checked against the exact artifacts if this project ever redistributes.

6. How Veronica.v4 uses this checkpoint

6.1 Serving stack (smoke / first-chat runtime)

Pinned in `config/runpod-core.json` and executed on successful starts:

- Image:
`runpod/pytorch@sha256:0a360022e8de4375af99430f84e8b38951acc397252163a37ceac7204d01be35` (tag `runpod/pytorch:1.0.2-cu1281-torch280-ubuntu2404`)
- vLLM 0.11.0, Transformers 4.57.1, isolated runtime under
`/workspace/veronica-core/runtime-vllm-0.11.0`
- GPU: 1x NVIDIA A100-SXM4-80GB, Secure Cloud, ceiling \$1.75/hour, observed \$1.59/hour
- Flags: `--dtype bfloat16 --max-model-len 8192 --max-num-seqs 1 --gpu-memory-utilization 0.92 --enforce-eager --served-model-name Veronica`
- First-chat load: 239.9 seconds, 56.9342 GiB model; idle allocation after serving 75,149 MiB including cache/runtime
- One streaming probe: 1.167 s to first content, 18.71 completion tokens per total second (smoke, not a benchmark)

T2 frozen comparison runtime is different: vLLM 0.28.0 and Transformers 5.8.0, context 32,768, because Candidate B's card requires those versions. The 0.11.0 smoke runtime is not accepted as the four-model comparison runtime.

6.2 Integrity pipeline

`scripts/prepare_runpod_model.py` downloads the pinned revision into `.uploading-<revision>`, verifies every expected path's byte count and SHA-256, then atomically renames staging to the canonical directory. A bad existing copy is never overwritten. Interrupted downloads remain in staging. That is the project's implementation of "validate transfers before promotion from `.uploading` storage."

6.3 Wrapper layer (not weight changes)

`src/veronica_core` injects:

- A system persona naming the assistant Veronica (the checked-in persona text currently addresses "TRAVIS"; the owner-facing product name is Raine — a wrapper inconsistency, not a weight edit).
- Mode presets: Chat, Deep Reasoning, Creative, Coding. SOURCE-OF-TRUTH: these are prompt presets until native model controls are qualified. Candidate A cannot be turned into a thinking model by a preset.
- Public alias rewrite, health/capabilities endpoints, streaming SSE, browser UI.

The persona also asserts an uncensored, non-refusing policy in prompt text. That is prompting, stacked on top of community ablation. It is not evidence that native refusal is gone, and it is not a fine-tune.

6.4 Timeline of project actions on this exact revision

When	What happened to Candidate A	What did not happen
2026-08-30	Provenance pin; first supervised A100 run; 27-file hash verification; first real API/UI chat; Pod terminated; volume retained	No fine-tune; no model selection
2026-08-31	Cold restart of the same weights; retained UI transcript; evaluation bank and recorded-conversation review	No training on the transcript
2026-09-01	T2 protocol freeze: A vs official control vs B vs B-control; license/lineage snapshots without weight downloads for the other three	No live T2 matrix
2026-09-04	Wrapper streaming, session persistence, Markdown, message controls	No GPU; no weight change
2026-09-06	Multiple authorized start attempts; some cancelled; some reached UI; stock/network issues documented	No Candidate B transfer
2026-09-08 T040904Z	Same revision re-verified and served; wrapper health ready at 05:33:08 UTC; Pod h33pceh6ug2q5x; deadline 06:12:36 UTC	Still not T2; still not selected

7. Observed behavior (smoke and human review, not T2)

These findings are about this checkpoint as wrapped and served, not official Qwen scores.

First-chat (2026-08-30) manual review
(runs/2026-08-30-supervised-first-chat/manual-review.md):

- Creative four-sentence scene: produced. Broader prose not qualified.
- Coding `is_even`: generated function executed locally and passed the model's three asserts plus six extra integer cases. The model also claimed assertions passed though no execution tool existed — action-truthfulness failure.
- Probability of two red balls from 3 red + 2 blue without replacement: opened with 3/5, then derived 3/10. Contradiction, not a clean reasoning pass. The 2026-09-08 wrapper-smoke reasoning item repeats the same 3/5-then-3/10 pattern.
- Direct provider invented autobiographical travel; wrapper identified as AI. Identity consistency unproven.
- Memory language ("always remember") exceeded actual session context. No long-term memory module exists.
- UI two-turn recall of "silver compass" worked for that example only.

Recorded conversation (2026-08-31) owner-adjudicated: human mean 0.667/4, two critical failures (invented calendar/email monitoring and learning; invented telemetry, voice, and timings after being told it was hallucinating). Gate: `blocked_on_observed_failures`.

2026-09-08 wrapper-smoke still advertises `reasoning_content: null` (consistent with non-thinking Instruct) and `basicSmokePassed: true` with the explicit caveat that semantic

correctness requires manual review. Capability qualification remains pending.

8. What has not been done

1. No LoRA, QLoRA, full fine-tune, merge, or quantization of these weights by Veronica.v4.
2. No hash-validated copy of the official Instruct-2507 control, Candidate B, or B-control on the volume.
3. No live T2 four-model, persona-free, matched-runtime comparison (vLLM 0.28.0 / Transformers 5.8.0 / 32,768 context).
4. No signed model-selection decision. Registry `selectionStatus` is `benchmark_required`.
5. No native tool-call parser qualification on this serve (hermes is reserved for the T2 A profile, not used on the 0.11.0 smoke command).
6. No long-context stress at 262K or 1M; configured context is 8,192.
7. No 48 GB quantization selection with measured VRAM. BF16 shards total ~61 GB on disk; the A100 80 GB smoke fit at ~57 GiB weights. A 48 GB GPU is not evidenced.
8. Wrapper "Deep Reasoning" is not native thinking.

9. Implications for the next milestone

Gate R ("Engine Accounted For") still wants a complete manifest for every T2 model, not only A; transfer validation exists in code and was proven for A; GPU/price/deadline are recorded (A100, \$1.75/hour ceiling, per-run duration). Gate T2 ("Mind Proven") cannot close on this paper or on smoke chats.

The honest next live experiment, when freshly authorized, is still the frozen T2 protocol: serve this exact revision and its official control, then B and B-control, on the matched 0.28.0 runtime, and score the 60-case pack plus executable coding, schema, tools, and long-context stress — without treating wrapper persona as intelligence.

Until that exists, Candidate A is a provisional, hash-verified, actually serving uncensored Instruct-2507 MoE with documented contradiction and action-truthfulness failures. It is the right installed model. It is not a finished Veronica mind.

10. Sources

Primary Hub and cards

1. <https://huggingface.co/huihui-ai/Huihui-Qwen3-30B-A3B-Instruct-2507-abliterated> — Huihui Candidate A card (fetched 2026-09-08; matches pinned README).
2. [runs/2026-09-08T040904Z-start-veronica/provenance/candidate/README.md](https://huggingface.co/runs/2026-09-08T040904Z-start-veronica/provenance/candidate/README.md) and [hub-metadata.json](https://huggingface.co/runs/2026-09-08T040904Z-start-veronica/provenance/candidate/hub-metadata.json) — pinned revision snapshot, including sha `e2f73ec7e99ee316beb8069ca90e4c3cbef8aa0f`.
3. <https://huggingface.co/Qwen/Qwen3-30B-A3B-Instruct-2507> — official Instruct-2507 card.
4. [runs/2026-09-01-t2-qualification/provenance/candidate-a-control/README.md](https://huggingface.co/runs/2026-09-01-t2-qualification/provenance/candidate-a-control/README.md) — pinned control card.
5. <https://github.com/QwenLM/Qwen3> — Qwen3-2507 release notes and Apache-2.0 statement.
6. Qwen Team, *Qwen3 Technical Report*, arXiv:2505.09388.

Abliteration

7. Arditi et al., *Refusal in Language Models Is Mediated by a Single Direction*, arXiv:2406.11717 (NeurIPS 2024).
8. <https://github.com/Sumandora/remove-refusals-with-transformers> — implementation Huihui cites.
9. https://github.com/andyrdt/refusal_direction — Arditi et al. code.
10. <https://github.com/Orion-zhen/abliteration> — community weight-bake formula and "not uncensorship" note.

License

11. <https://www.apache.org/licenses/LICENSE-2.0>
12. [runs/2026-09-01-t2-qualification/license-review.md](#)

Project evidence (this exact revision)

13. [config/model-registry.json](#), [config/runpod-core.json](#), [config/t2-qualification.json](#)
14. [docs/SOURCE-OF-TRUTH.md](#), [TODO.md](#), [AGENTS.md](#)
15. [runs/2026-08-30-supervised-first-chat/decision.md](#), [manual-review.md](#), [validated-model-manifest.json](#)
16. [runs/2026-09-08T040904Z-start-veronica/](#) — [validated-model-manifest.json](#), [expected-model-manifest.json](#), [server-command.json](#), [volume-inspection.txt](#), [wrapper-smoke.json](#), [profile.json](#)
17. [runs/2026-09-01-t2-qualification/decision.md](#)
18. [runs/2026-08-31-recorded-conversation-eval/report.md](#)
19. [scripts/prepare_runpod_model.py](#)
20. [src/veronica_core/persona.py](#), [src/veronica_core/app.py](#)

11. Independent agent cross-check (2026-09-08)

A bounded deep-research workflow (deep-research-2) independently verified Candidate A identity, Apache-2.0, 13-shard BF16 dump, volume v53gj9flzs, vLLM 0.11.0 serve alias Veronica, absence of Candidate B on the volume listing, first-chat closeout, recorded-conversation gate, and T2 still frozen. Raw agent report: 2026-09-08-candidate-a-deep-research-2-agent-report.md. Status: Partial.

Rejected mix-up: some researchers inspected `/home/dubs/projects/veronica-v4-LOCAL` and local Ollama tag `Veronica:latest` (a bartowski dolphin-2.9-llama3-8b Q4_K_M GGUF, ~4.92 GB). That leftover tree is not canonical Veronica.v4. Canonical [docs/SOURCE-OF-TRUTH.md](#) mentions Ollama only as an analogy for an alias/Modelfile; this project serves Candidate A through vLLM/RunPod. Do not treat Dolphin 3 / `veronica-v.2:dolphin3-8b-uncensored` as the installed foundation.

Adopted extras from that run: official Qwen 16-shard vs Huihui 13-shard packing; Hub commit titles (Huihui `e2f73ec7` "Upload Modelfile" 2025-08-02; Qwen control `0d7cf239` "Update tokenizer_config.json" 2025-09-17); Sumandora public code is activation ablation without a weight exporter; expected-manifest SHA-256 vs git-blob split documented above.

12. Coverage and uncertainty

- Huihui's "new and faster method" is not specified at the layer/tensor level for Candidate A. Independent reconstruction of the ablation is not claimed. The published Sumandora recipe ablates activations at inference; Huihui's artifact is a weight dump whose exact bake recipe is unpublished.
- Official Qwen benchmarks are for the untouched Instruct-2507, not this derivative, and not this project's 8k-context vLLM 0.11.0 serve.
- Apache-2.0 documentation is not a lawyer's opinion.
- Live Pod presence after 2026-09-08T06:12:36Z is not asserted; health ready was observed at 05:33:08Z that day.
- This paper does not authorize paid compute, training, or foundation selection.